

以下给出了每个练习题对应的代码部分和相应的注释，其中一些注释是在报告中添加的以更好地解释实现过程。

练习题 1: 完成 `kernel/mm/buddy.c` 中的 `split_chunk`、`merge_chunk`、`buddy_get_pages`、和 `buddy_free_pages` 函数中的 `LAB 2 TODO 1` 部分，其中 `buddy_get_pages` 用于分配指定阶大小的连续物理页，`buddy_free_pages` 用于释放已分配连续物理页。

```
static struct page *split_chunk(struct phys_mem_pool *pool, int order,
                                struct page *chunk)
{
    while (chunk -> order > order){ // Before reaching to the aimed oerder
        struct page *buddy;
        -- chunk -> order; // decrease the order after splitting
        buddy = get_buddy_chunk(pool, chunk);
        buddy -> order = chunk -> order;
        buddy -> allocated = 0; // allocated is false
        ++ pool -> free_lists[buddy -> order].nr_free; // increase the number of free
        list_add(&(buddy -> node), &(pool -> free_lists[buddy -> order].free_list)); // add the buddy into free list of
        corresponding order
    }
    return chunk;
}

static struct page *merge_chunk(struct phys_mem_pool *pool, struct page *chunk)
{
    while(chunk -> order < BUDDY_MAX_ORDER - 1){
        struct page *buddy = get_buddy_chunk(pool, chunk);
        if (buddy == NULL || buddy -> allocated || buddy -> order != chunk -> order){
            break;
        } // no buddy available or wrong buddy
        -- pool -> free_lists[chunk -> order].nr_free; // decrease free number of corresponding order after merge

        list_del(&(buddy -> node));

        if (chunk > buddy){
            chunk = buddy;
        } // make chunk the one with lower addr

        ++ chunk -> order;
        chunk -> allocated = 1;
    }
    return chunk;
}

struct page *buddy_get_pages(struct phys_mem_pool *pool, int order)
{
    int cur_order;
    struct list_head *free_list;
    struct page *page = NULL;

    if (unlikely(order >= BUDDY_MAX_ORDER)) {
        kwarn("ChCore does not support allocating such too large "
            "contious physical memory\n");
        return NULL;
    }

    lock(&pool->buddy_lock);

    /* BLANK BEGIN */
    for (cur_order = order; cur_order < BUDDY_MAX_ORDER; ++ cur_order){
        if(pool -> free_lists[cur_order].nr_free > 0){
            free_list = pool -> free_lists[cur_order].free_list.next; // first one is head */
            page = list_entry(free_list, struct page, node); // get the addr of the page which has free_list's addr as
            node*/
            -- pool -> free_lists[page -> order].nr_free;
            list_del(&(page -> node));
            break;
        }
    }
    if(page == NULL) {
        unlock(&pool -> buddy_lock);
        return NULL;
    }
    page = split_chunk(pool, order, page);
    page -> allocated = 1;
    /* BLANK END */
}
```

out:

```

        unlock(&pool->buddy_lock);
        return page;
    }

void buddy_free_pages(struct phys_mem_pool *pool, struct page *page)
{
    int order;
    struct list_head *free_list;

    lock(&pool->buddy_lock);

    /* BLANK BEGIN */
    if (page == NULL || page -> allocated == 0) {return;}
    page -> allocated = 0; // free the page
    struct page *merged_chunk = merge_chunk(pool, page);
    merged_chunk -> allocated = 0;
    order = merged_chunk -> order;
    ++ pool -> free_lists[order].nr_free;
    free_list = &(pool -> free_lists[order].free_list);
    list_add(&(merged_chunk -> node), free_list);
    /* BLANK END */

    unlock(&pool->buddy_lock);
}

```

练习题 2: 完成 `kernel/mm/slab.c` 中的 `choose_new_current_slab`、`alloc_in_slab_impl` 和 `free_in_slab` 函数中的 LAB 2 TODO 2 部分, 其中 `alloc_in_slab_impl` 用于在 slab 分配器中分配指定阶大小的内存, 而 `free_in_slab` 则用于释放上述已分配的内存。

```

static void choose_new_current_slab(struct slab_pointer *pool, int order)
{
    /* BLANK BEGIN */
    if (list_empty(&pool -> partial_slab_list)){
        pool -> current_slab = init_slab_cache(order, SIZE_OF_ONE_SLAB); // no available slab, allocate new one
    } else {
        pool -> current_slab = list_entry(pool -> partial_slab_list.next, struct slab_header, node); /* choose the
head->next in the partial_slab_list */
        list_del(pool -> partial_slab_list.next);
    }

    /* BLANK END */
}

static void *alloc_in_slab_impl(int order)
{
    struct slab_header *current_slab;
    struct slab_slot_list *free_list;
    void *next_slot;

    lock(&slabs_locks[order]);

    current_slab = slab_pool[order].current_slab;
    /* When serving the first allocation request. */
    /* unlikely is used to optimize branch, which tells compiler this case is unlikely to happen */
    if (unlikely(current_slab == NULL)) {
        current_slab = init_slab_cache(order, SIZE_OF_ONE_SLAB);
        if (current_slab == NULL) {
            unlock(&slabs_locks[order]);
            return NULL;
        }
        slab_pool[order].current_slab = current_slab;
    }

    /* BLANK BEGIN */
    if (current_slab->current_free_cnt == 0) { // no available slot!
        choose_new_current_slab(&slab_pool[order], order);
    }
    free_list = current_slab->free_list_head;
    current_slab->free_list_head = ((struct slab_slot_list *) (current_slab->free_list_head))->next_free;
    // free_list->next_free = NULL;
    --current_slab->current_free_cnt; // use one slot

    /* BLANK END */
}

```

```

        unlock(&slabs_locks[order]);

        return (void *)free_list;
}

void free_in_slab(void *addr)
{
    struct page *page;
    struct slab_header *slab;
    struct slab_slot_list *slot;
    int order;

    slot = (struct slab_slot_list *)addr;
    page = virt_to_page(addr);
    if (!page) {
        kdebug("invalid page in %s", __func__);
        return;
    }

    slab = page->slab;
    order = slab->order;
    lock(&slabs_locks[order]);

    try_insert_full_slab_to_partial(slab);

#ifdef ENABLE_DETECTING_DOUBLE_FREE_IN_SLAB == ON
    /*
     * SLAB double free detection: check whether the slot to free is
     * already in the free list.
     */
    if (check_slot_is_free(slab, slot) == 1) {
        kinfo("SLAB: double free detected. Address is %p\n",
              (unsigned long)slot);
        BUG_ON(1);
    }
#endif

    /* BLANK BEGIN */
    slot->next_free = slab->free_list_head; // do link list operation to free the slot
    slab->free_list_head = (void *)slot;
    ++slab->current_free_cnt;
    /* BLANK END */

    try_return_slab_to_buddy(slab, order);

    unlock(&slabs_locks[order]);
}

```

练习题 3: 完成 `kernel/mm/kmalloc.c` 中的 `_kmalloc` 函数中的 **LAB 2 TODO 3** 部分, 在适当位置调用对应的函数, 实现 `kmalloc` 功能

```

void *_kmalloc(size_t size, bool is_record, size_t *real_size)
{
    void *addr;
    int order;

    if (unlikely(size == 0))
        return ZERO_SIZE_PTR;

    if (size <= SLAB_MAX_SIZE) {
        /* Step 1: Allocate in slab for small requests. */
        /* BLANK BEGIN */
        addr = alloc_in_slab(size, real_size);
        /* BLANK END */
#ifdef ENABLE_MEMORY_USAGE_COLLECTING == ON
        if (is_record && collecting_switch) {
            record_mem_usage(*real_size, addr);
        }
#endif
    } else {
        /* Step 2: Allocate in buddy for large requests. */
        /* BLANK BEGIN */
        order = size_to_page_order(size);
        addr = get_pages(order);
    }
}

```

```

        /* BLANK END */
    }

    BUG_ON(!addr);
    return addr;
}

```

练习题 4: 完成 `kernel/arch/aarch64/mm/page_table.c` 中的

`query_in_pgtbl`、`map_range_in_pgtbl_common`、`unmap_range_in_pgtbl` 和 `mprotect_in_pgtbl` 函数中的 LAB 2 TODO 4 部分，分别实现页表查询、映射、取消映射和修改页表权限的操作，以 4KB 页为粒度。

```

int query_in_pgtbl(void *pgtbl, vaddr_t va, paddr_t *pa, pte_t **entry)
{
    /* BLANK BEGIN */

    ptp_t *cur_ptp;
    pte_t *pte;
    int res;

    cur_ptp = (ptp_t *) pgtbl;
    /* L0 */
    res = get_next_ptp(cur_ptp, 0, va, &cur_ptp, &pte, false, NULL);
    if (res == -ENOMAPPING) {
        return res;
    }

    res = get_next_ptp(cur_ptp, 1, va, &cur_ptp, &pte, false, NULL);
    if (res == -ENOMAPPING) {
        return res;
    } else if (res == BLOCK_PTP) {
        if (entry != NULL){
            *entry = pte;
        }

        *pa = virt_to_phys(cur_ptp) + GET_VA_OFFSET_L1(va);
        return 0;
    }

    /* L2 */
    res = get_next_ptp(cur_ptp, 2, va, &cur_ptp, &pte, false, NULL);
    if (res == -ENOMAPPING) {
        return res;
    } else if (res == BLOCK_PTP) {

        if (entry != NULL){
            *entry = pte;
        }

        *pa = virt_to_phys(cur_ptp) + GET_VA_OFFSET_L2(va);
        return 0;
    }

    /* L3 */
    res = get_next_ptp(cur_ptp, 3, va, &cur_ptp, &pte, false, NULL);
    if (res == -ENOMAPPING) {
        return res;
    }
    if (entry != NULL){
        *entry = pte;
    }

    *pa = virt_to_phys(cur_ptp) + GET_VA_OFFSET_L3(va);
    return 0;

    /* BLANK END */

    return 0;
}

static int map_range_in_pgtbl_common(void *pgtbl, vaddr_t va, paddr_t pa, size_t len,
                                     vmr_prop_t flags, int kind, long *rss)
{
    /* BLANK BEGIN */
    u64 page_num = len / PAGE_SIZE + (len % PAGE_SIZE > 0); /* ROUND_UP */

```

```

while(page_num > 0){
    ptp_t *cur_ptp = (ptp_t *)pgtbl;
    pte_t *pte;
    for(int i = 0; i < 3; ++ i){
        get_next_ptp(cur_ptp, i, va, &cur_ptp, &pte, true, NULL);
    }

    for(int i = GET_L3_INDEX(va); i < PTP_ENTRIES; ++ i){
        pte_t new_pte;
        new_pte.pte = 0;
        new_pte.l3_page.is_valid = 1;
        new_pte.l3_page.is_page = 1;
        new_pte.l3_page.pfn = pa >> PAGE_SHIFT; /* page frame number */
        set_pte_flags(&new_pte, flags, kind);
        cur_ptp->ent[i] = new_pte;

        -- page_num;

        if(page_num == 0){
            break;
        }

        va += PAGE_SIZE; /* move on to next mapping */
        pa += PAGE_SIZE;
    }

}

/* BLANK END */
return 0;
}

int unmap_range_in_pgtbl(void *pgtbl, vaddr_t va, size_t len, long *rss)
{
    /* BLANK BEGIN */
    u64 page_num = len / PAGE_SIZE + (len % PAGE_SIZE > 0); /* ROUND_UP */
    while (page_num > 0) {
        ptp_t *cur_ptp = (ptp_t *)pgtbl;
        pte_t *pte;
        int res;
        // L0
        res = get_next_ptp(cur_ptp, 0, va, &cur_ptp, &pte, false, NULL);
        if (res == -ENOMAPPING) {
            page_num -= LO_PER_ENTRY_PAGES;
            va += LO_PER_ENTRY_PAGES * PAGE_SIZE;
            continue;
        }
        // L1
        res = get_next_ptp(cur_ptp, 1, va, &cur_ptp, &pte, false, NULL);
        if (res == -ENOMAPPING) {
            page_num -= L1_PER_ENTRY_PAGES;
            va += L1_PER_ENTRY_PAGES * PAGE_SIZE;
            continue;
        }
        // L2
        res = get_next_ptp(cur_ptp, 2, va, &cur_ptp, &pte, false, NULL);
        if (res == -ENOMAPPING) {
            page_num -= L2_PER_ENTRY_PAGES;
            va += L2_PER_ENTRY_PAGES * PAGE_SIZE;
            continue;
        }
        // L3
        for (int i = GET_L3_INDEX(va); i < PTP_ENTRIES; ++i) {
            cur_ptp->ent[i].pte = PTE_DESCRIPTOR_INVALID;

            va += PAGE_SIZE;

            --page_num;
            if (page_num == 0) {
                break;
            }
        }
    }
}

```

```

        /* BLANK END */

        dsb(ishst);
        isb();

        return 0;
}

int mprotect_in_pgtbl(void *pgtbl, vaddr_t va, size_t len, vmr_prop_t flags)
{
    /* BLANK BEGIN */
    u64 page_num = len / PAGE_SIZE + (len % PAGE_SIZE > 0); /* ROUND_UP */
    while(page_num > 0){
        ptp_t *cur_ptp = (ptp_t *)pgtbl;
        pte_t *pte;
        for(int i = 0; i < 3; ++ i){
            get_next_ptp(cur_ptp, i, va, &cur_ptp, &pte, true, NULL);
        }

        for(int i = GET_L3_INDEX(va); i < PTP_ENTRIES; ++ i){

            set_pte_flags(&(cur_ptp -> ent[i]), flags, USER_PTE); // Set flags

            va += PAGE_SIZE; /* move on to next mapping */

            -- page_num;

            if(page_num == 0){
                break;
            }
        }

        /* BLANK END */
        return 0;
    }
}

```

思考题 5: 阅读 Arm Architecture Reference Manual, 思考要在操作系统中支持写时拷贝 (Copy-on-Write, CoW) 需要配置页表描述符的哪个/哪些字段, 并在发生页错误时如何处理。(在完成第三部分后, 你也可以阅读页错误处理的相关代码, 观察 ChCore 是如何支持 Cow 的)

L3页表项中的 AP 字段用于定义物理页的读写权限。实现写时拷贝机制时, 该物理页应设为可读不可写, 因此 AP 字段应设置为 11。当有应用程序试图对这个物理页进行写操作时, 会引发一个访问权限异常。随后, 操作系统会复制这个物理页, 将复制页的 AP 字段设置为 01 (即可读可写), 更新相应的页表项, 然后让应用程序继续执行写操作。

思考题 6: 为了简单起见, 在 ChCore 实验 Lab1 中没有为内核页表使用细粒度的映射, 而是直接沿用了启动时的粗粒度页表, 请思考这样做有什么问题。

粗粒度页表一次映射2M的页, 可能会在后续使用过程中产生很多内部碎片。

练习题 8: 完成 `kernel/arch/aarch64/irq/pgfault.c` 中的 `do_page_fault` 函数中的 LAB 2 TODO 5 部分, 将缺页异常转发给 `handle_trans_fault` 函数。

在填空部分填入以下一行内容即可。

```
ret = handle_trans_fault(current_thread -> vmSPACE, fault_addr);
```

练习题 9: 完成 `kernel/mm/vmSPACE.c` 中的 `find_vmr_for_va` 函数中的 LAB 2 TODO 6 部分, 找到一个虚拟地址找在其虚拟地址空间中的 VMR。

```

struct vmregion *find_vmr_for_va(struct vmSPACE *vmSPACE, vaddr_t addr)
{
    /* BLANK BEGIN */
    // search the vmr which includes this va
    struct rb_node *node = rb_search(&(vmSPACE->vmr_tree), (const void*)addr, cmp_vmr_and_va);
    if (node == NULL) {
        return NULL; // If not found
    }
    struct vmregion *res = rb_entry(node, struct vmregion, tree_node); // get the vmr that has node as its
    `tree_node`
    return res;
    /* BLANK END */
}

```

练习题 10: 完成 `kernel/mm/pgfault_handler.c` 中的 `handle_trans_fault` 函数中的 `LAB 2 TODO 7` 部分（函数内共有 3 处填空，不要遗漏），实现 `PMO_SHM` 和 `PMO_ANONYM` 的按需物理页分配。你可以阅读代码注释，调用你之前见到过的相关函数来实现功能。

```

int handle_trans_fault(struct vmSPACE *vmSPACE, vaddr_t fault_addr)
{
    struct vmregion *vmr;
    struct pmoobject *pmo;
    paddr_t pa;
    unsigned long offset;
    unsigned long index;
    int ret = 0;

    lock(&vmSPACE->vmSPACE_lock);
    vmr = find_vmr_for_va(vmSPACE, fault_addr);

    if (vmr == NULL) {
        kinfo("handle_trans_fault: no vmr found for va 0x%lx!\n",
            fault_addr);
        dump_pgfault_error();
        unlock(&vmSPACE->vmSPACE_lock);

#ifdef CHCORE_ARCH_AARCH64 || defined(CHCORE_ARCH_SPARC)
        /* kernel fault fixup is only supported on AArch64 and Sparc */
        return -EFAULT;
#endif

        sys_exit_group(-1);

        BUG("should not reach here");
    }

    pmo = vmr->pmo;
    /* Get the offset in the pmo for faulting addr */
    offset = ROUND_DOWN(fault_addr, PAGE_SIZE) - vmr->start + vmr->offset;
    vmr_prop_t perm = vmr->perm;
    switch (pmo->type) {
    case PMO_ANONYM:
    case PMO_SHM: {
        /* Boundary check */
        BUG_ON(offset >= pmo->size);

        /* Get the index in the pmo radix for faulting addr */
        index = offset / PAGE_SIZE;

        fault_addr = ROUND_DOWN(fault_addr, PAGE_SIZE);

        pa = get_page_from_pmo(pmo, index);
        if (pa == 0) {

            long rss = 0;

            /* BLANK BEGIN */

            pa = virt_to_phys((vaddr_t)get_pages(0)); // allocate a physical page
            memset((void*)phys_to_virt(pa), 0, PAGE_SIZE); // clear it to 0
            /* BLANK END */

            kdebug("commit: index: %ld, 0x%lx\n", index, pa);
            commit_page_to_pmo(pmo, index, pa);

            /* Add mapping in the page table */
            lock(&vmSPACE->pgtbl_lock);

```

```

        /* BLANK BEGIN */
        map_range_in_pgtbl(vmspace -> pgtbl, fault_addr, pa, PAGE_SIZE, perm, &rss);
        /* BLANK END */
        vmspace->rss += rss;
        unlock(&vmspace->pgtbl_lock);
    } else {
        if (pmo->type == PMO_SHM || pmo->type == PMO_ANONYM) {
            /* Add mapping in the page table */
            long rss = 0;
            lock(&vmspace->pgtbl_lock);
            /* BLANK BEGIN */
            map_range_in_pgtbl(vmspace -> pgtbl, fault_addr, pa, PAGE_SIZE, perm, &rss);
            /* BLANK END */
            /* LAB 2 TODO 7 END */
            vmspace->rss += rss;
            unlock(&vmspace->pgtbl_lock);
        }
    }

    if (perm & VMR_EXEC) {
        arch_flush_cache(fault_addr, PAGE_SIZE, SYNC_IDCACHE);
    }

    break;
}

case PMO_FILE: {
    unlock(&vmspace->vmspace_lock);
    fault_addr = ROUND_DOWN(fault_addr, PAGE_SIZE);
    handle_user_fault(pmo, ROUND_DOWN(fault_addr, PAGE_SIZE));
    BUG("Should never be here!\n");
    break;
}

case PMO_FORBID: {
    kinfo("Forbidden memory access (pmo->type is PMO_FORBID).\n");
    dump_pgfault_error();

    unlock(&vmspace->vmspace_lock);
    sys_exit_group(-1);

    BUG("should not reach here");
    break;
}

default: {
    kinfo("handle_trans_fault: faulting vmr->pmo->type"
        "(pmo type %d at 0x%lx)\n",
        vmr->pmo->type,
        fault_addr);
    dump_pgfault_error();

    unlock(&vmspace->vmspace_lock);
    sys_exit_group(-1);

    BUG("should not reach here");
    break;
}
}

unlock(&vmspace->vmspace_lock);
return ret;
}

```